



SocioScope - Interim Presentation

Ofir Fichman, Pavel Fadeev, Israel Peled



Project Review

- Developing a real-time "Public Pulse" system that monitors Twitter to quantify public sentiment across 10 key socio-economic pillars.
- Transitioned from manual annotation to an automated Synthetic Data Pipeline using LLMs to generate a high-quality, 10,000-sample dataset.
- Applying advanced ABSA to the socio-economic domain (beyond standard product reviews) by outputting a 10-dimensional sentiment vector per tweet.

Why Use SocioScope?



Nuanced Understanding:

Moving beyond simple 'Positive/Negative' labels to capture complex public opinion.

For Example:

A citizen might support the Government (+1) while simultaneously criticizing the Cost of Living (-1) in the same tweet.



Empowering Policymakers:

Providing mayors and government officials with precise analytics to understand exactly what troubles the public, enabling data-driven decision-making.

Previous work

Title / Year	The Task	Methods	Data	Results	Relation to SocioScope
<u>MEMD-ABSA: A Multi-Element Multi-Domain Dataset for Aspect-Based Sentiment Analysis (2023)</u>	Extracting the Aspect, Category, Opinion, and Sentiment even when they are not explicitly mentioned.	Generative Baselines (BART/T5) + Multi-Domain Training	20,000 sentences across 5 domains (Books, Clothing, etc.)	Revealed that mining implicit aspects and opinions remains the biggest challenge in open-domain ABSA.	Directly validates our use of LLMs to infer implicit socio-economic sentiment from vague tweets.
<u>Aspect-Based Sentiment Analysis Using BERT(2019)</u>	Advanced Modeling: Using pre-trained BERT to identify aspects and sentiments	Fine-tuned BERT + Sentence-Pair Modeling	SemEval-2015 & SemEval-2016	Outperformed previous SVM/CRF baselines	Justifies our use of BERT-like Transformers and "Supervised" synthetic labels.
<u>Sentiment Analysis in the Era of LLMs: A Reality Check(2023)</u>	Evaluating LLMs vs. Small Models	Comparing Zero-shot LLMs (ChatGPT) vs. Fine-tuned models (BERT)	26 Standard sentiment datasets (IMDB, Twitter, Rest14)	Fine-tuned small models (like BERT) often outperform Zero-shot LLMs in specific tasks.	Supports our hypothesis that a fine-tuned BERT can surpass Gemma's zero-shot performance.

Data Generation Pipeline



1. Define Attributes

Defined a list of aspects, personas, and writing styles for the model to generate tweets from.

2. Generate Target Vector

Created a function that generates a target vector, selecting 1-2 active aspects and assigning them non-zero sentiment.

3. Construct Dynamic Prompts

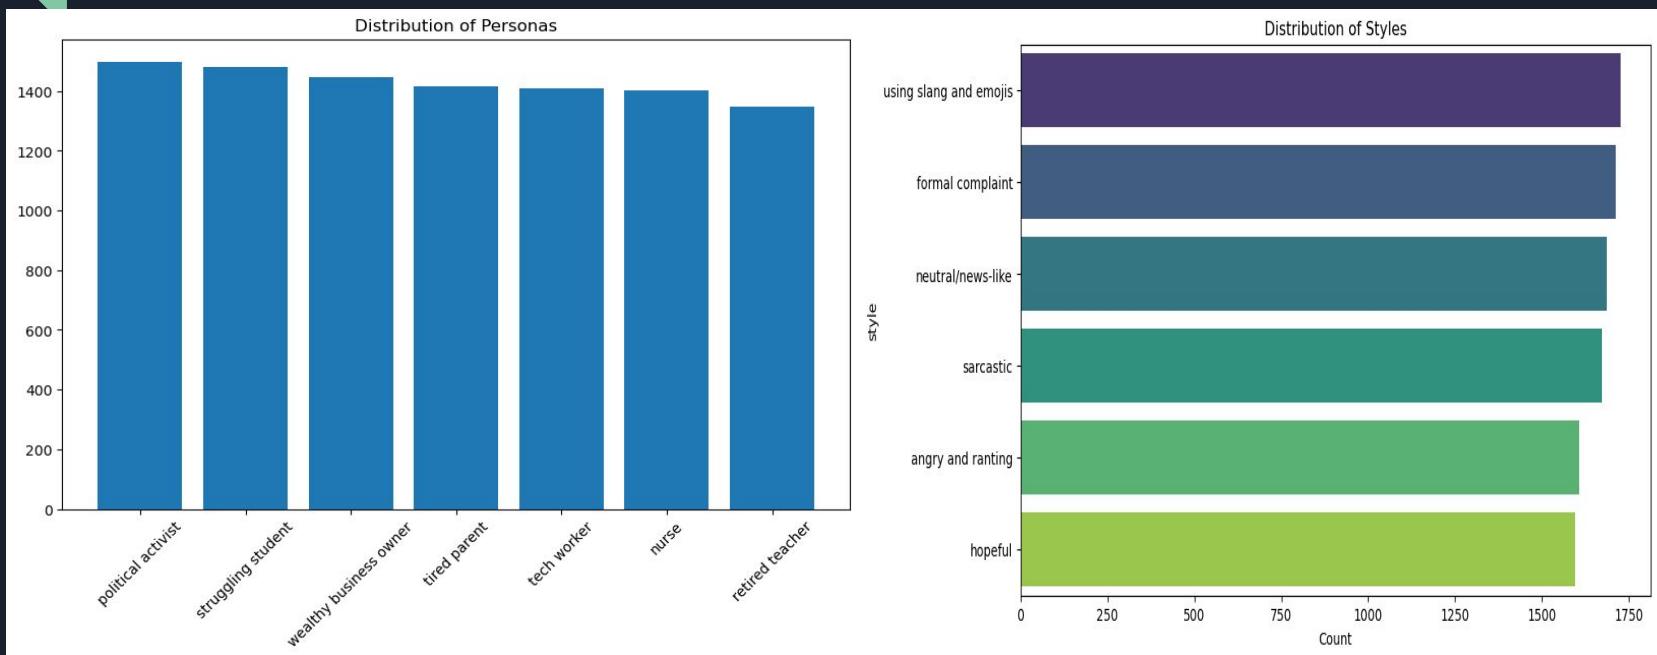
Constructed dynamic prompts that instruct the model to adopt a persona and write about the active aspects in a specific style.

4. Main Loop & LLM Query

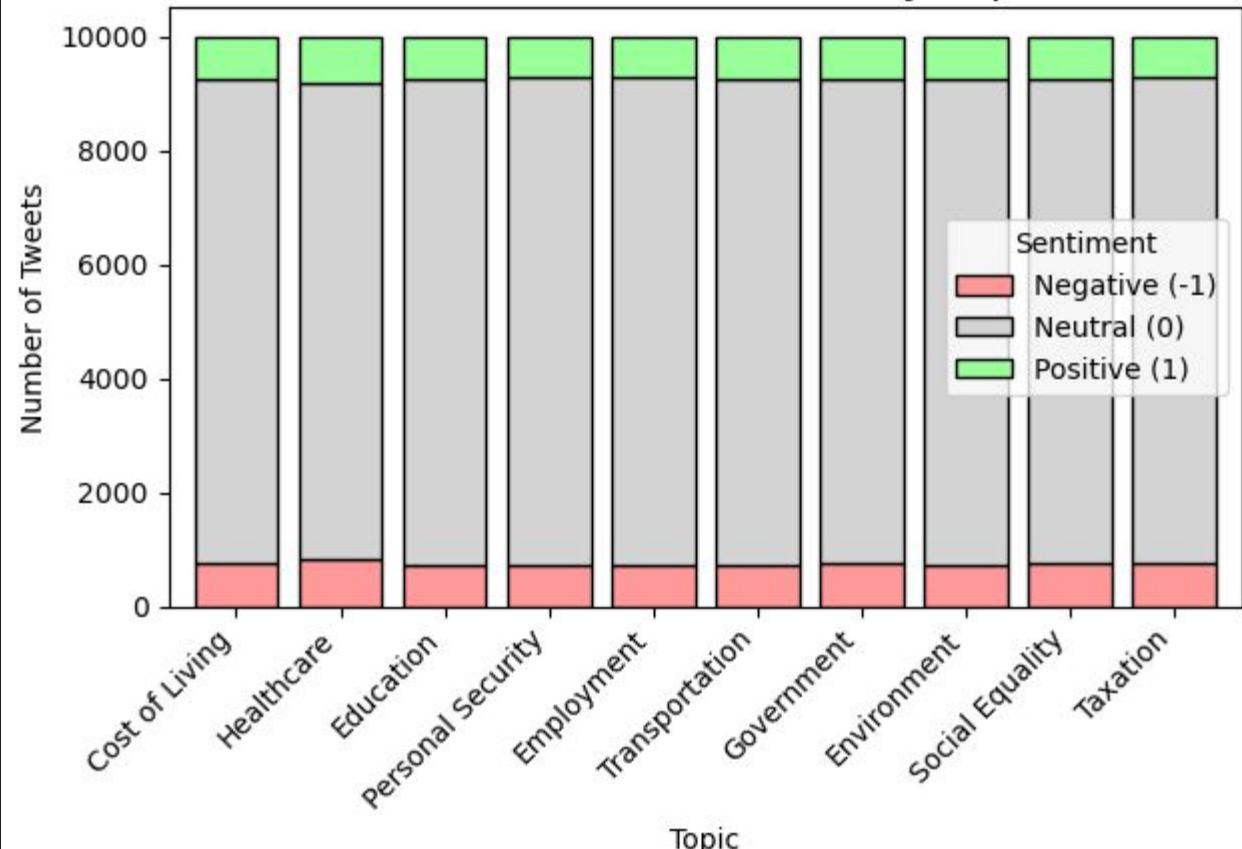
Implemented the main loop to orchestrate the pipeline: it randomizes attributes, queries the LLM via Ollama, and aggregates the results.

Dataset

Distributions

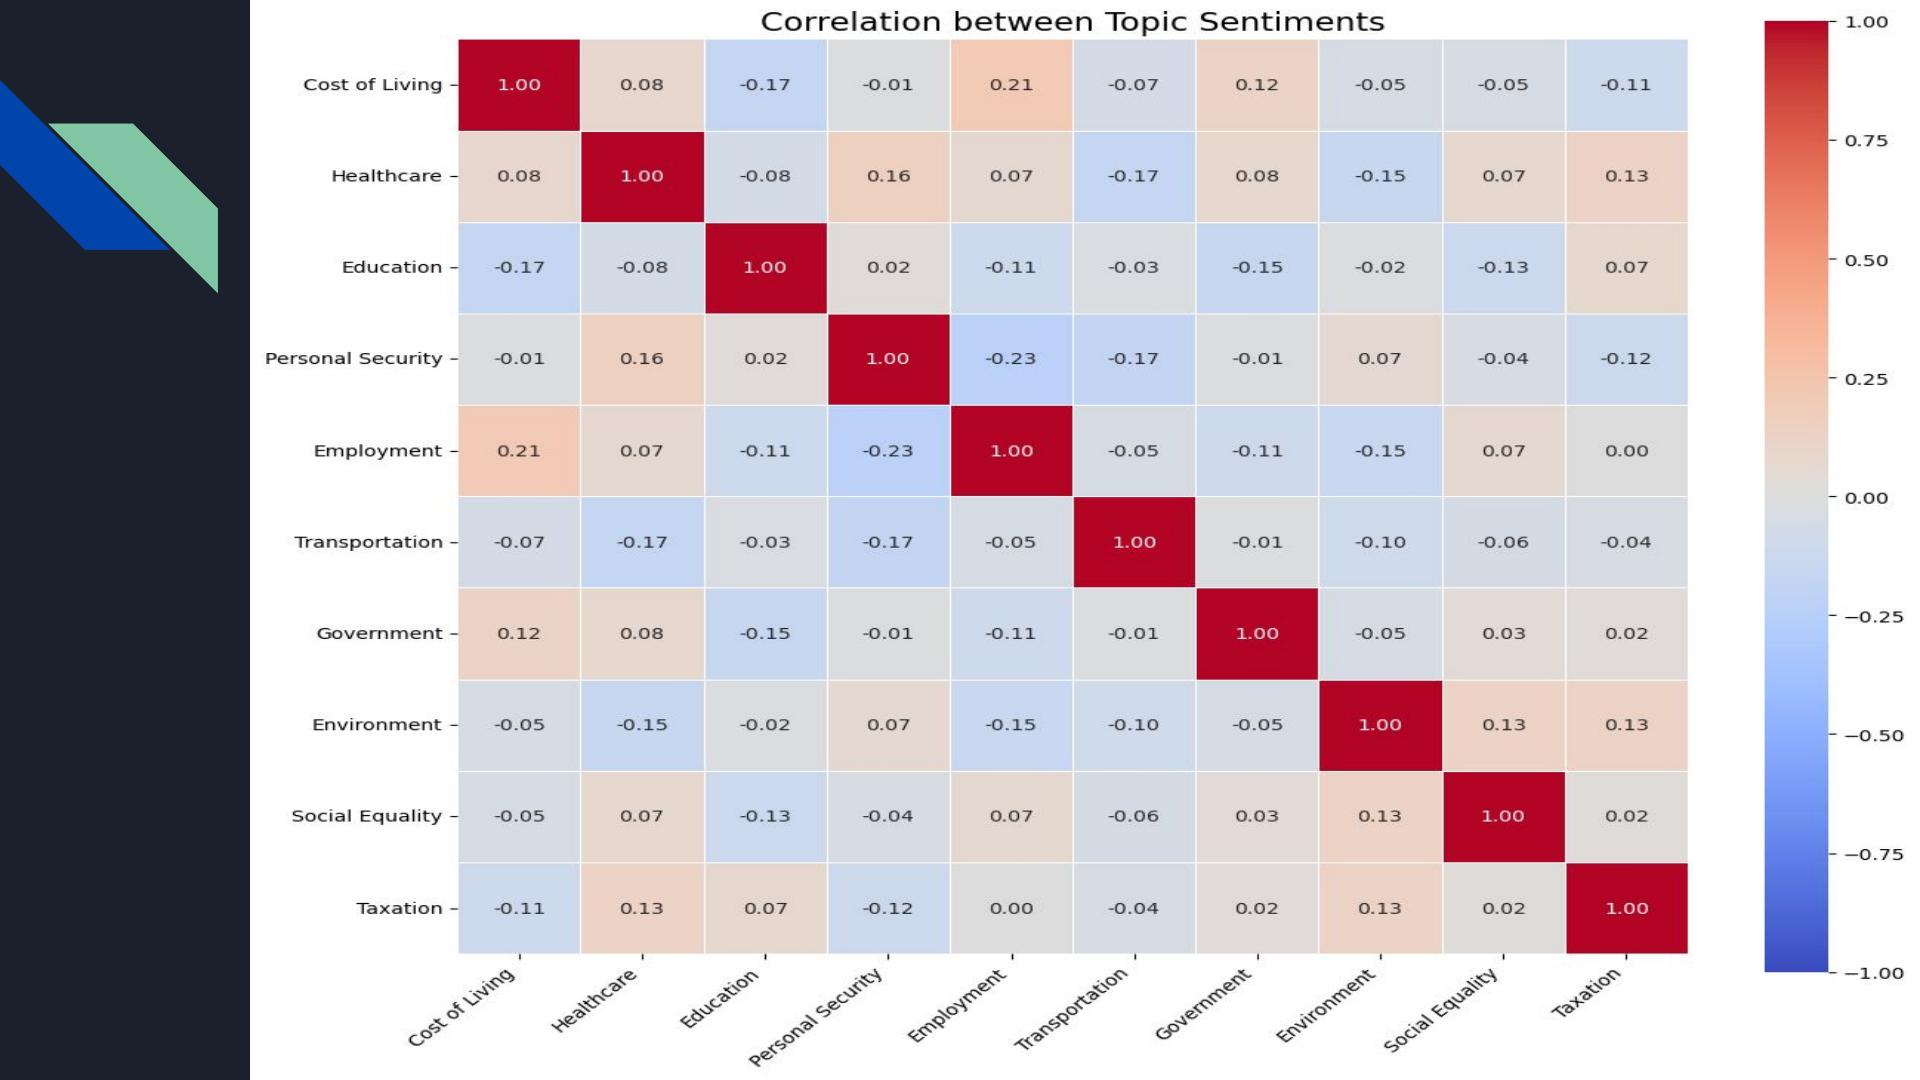


Sentiment Distribution by Topic

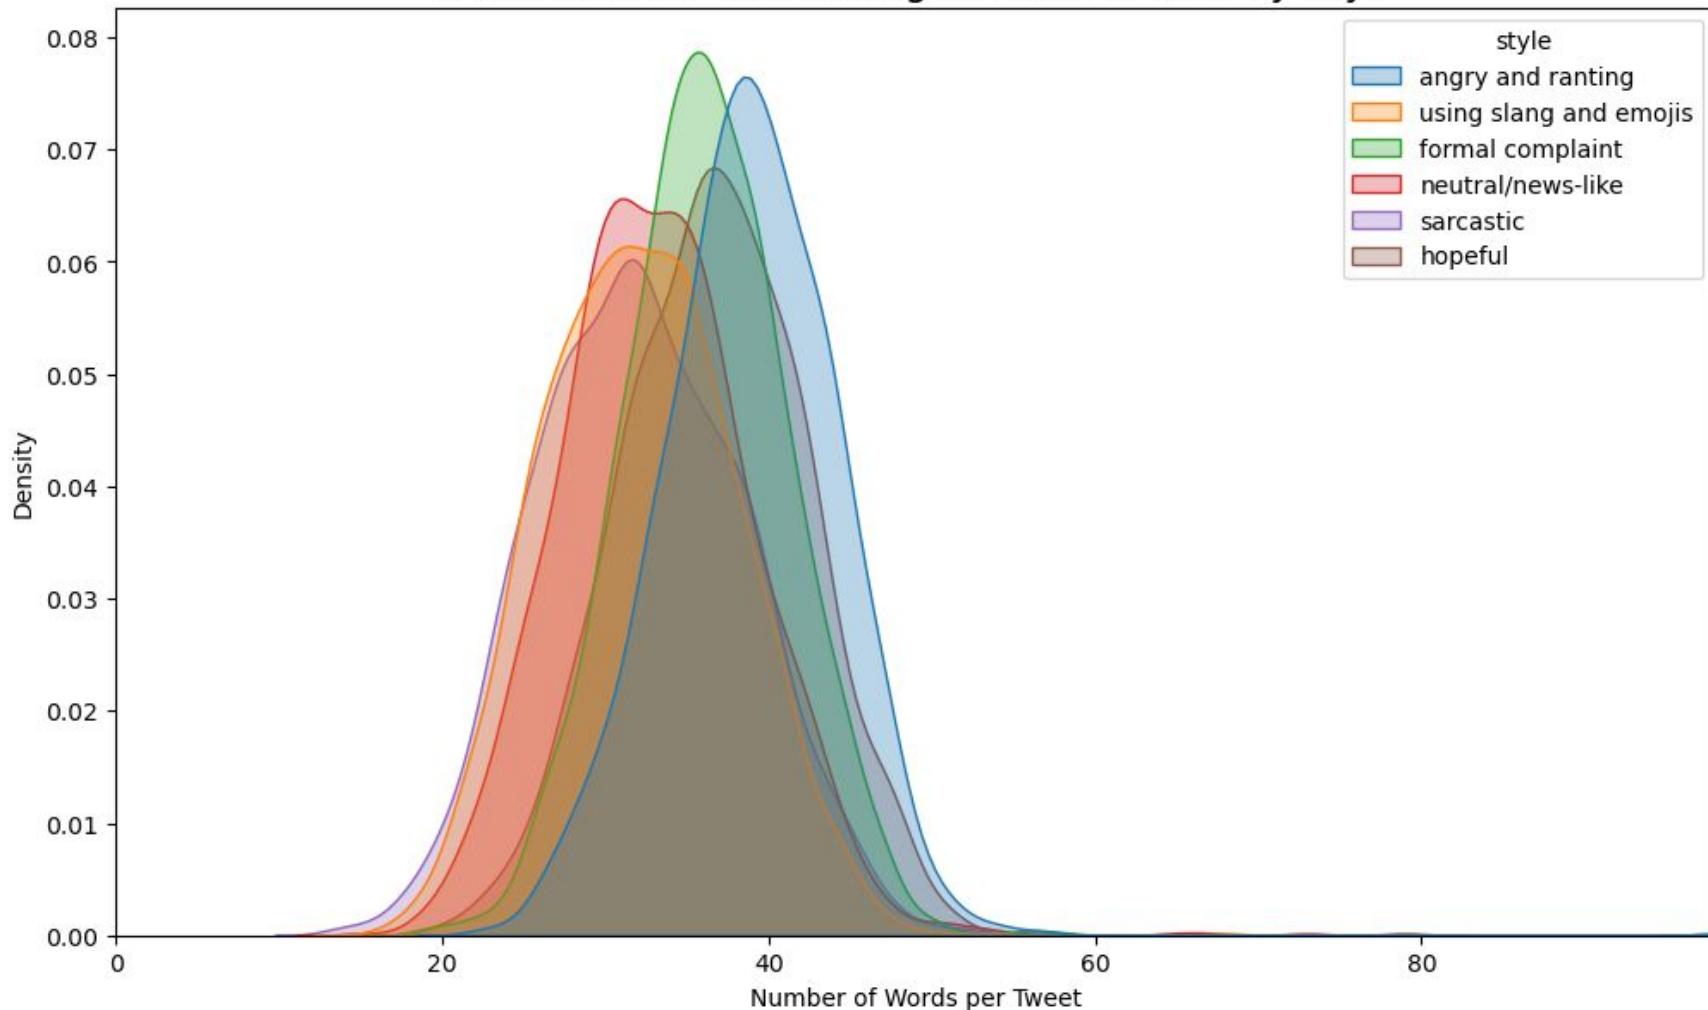


Average Sentiment (Excluding Neutrals) by Persona





Distribution of Tweet Length (Word Count) by Style



Baseline

```
ASPECTS = [
    "Cost of Living", "Healthcare", "Education", "Personal Security",
    "Employment", "Transportation", "Government", "Environment",
    "Social Equality", "Taxation"
]

def analyze_tweet_sentiment(tweet_text, model="gemma3"):
    prompt = f"""
    You are a precise data labeling assistant.
    Analyze the sentiment of the following tweet regarding these specific aspects:
    {ASPECTS}

    For each aspect, assign one of the following scores:
    1 : Positive sentiment
    -1 : Negative sentiment
    0 : Neutral sentiment OR the aspect is not mentioned in the tweet.

    Tweet: "{tweet_text}"

    Output Format:
    Return ONLY a raw JSON object with the aspects as keys and the scores (integer) as values.
    Do not write any introduction or explanation.
    """

    try:
        response = ollama.chat(model=model, messages=[
            {'role': 'user', 'content': prompt}
        ], format='json')

        content = response['message'][0]['content']
        result_dict = json.loads(content)
        final_vector = {aspect: result_dict.get(aspect, 0) for aspect in ASPECTS}

        return final_vector
    except Exception as e:
        print(f"Error processing tweet: {e}")

    return {aspect: 0 for aspect in ASPECTS}

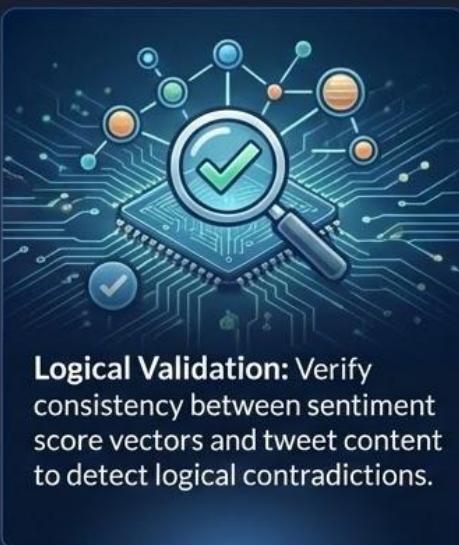
sample_tweet = "The air quality in this city is terrible because of the factories, but at least the new train system is fast and cheap."
sentiment_vector = analyze_tweet_sentiment(sample_tweet)
print("Tweet:", sample_tweet)
print("\nSentiment Vector:")
print(json.dumps(sentiment_vector, indent=4))
```

```
from tqdm import tqdm
tqdm.pandas()
sentiment_results
sampled_df['TweetText'].progress_apply(lambda x: analyze_tweet_sentiment(x))
final_df = pd.concat([sampled_df, sentiment_df], axis=1)
```

Comparison DataFrame created!
Average Model Accuracy: 82.60%

Is our work done? Absolutely not.

Current Objectives:



Logical Validation: Verify consistency between sentiment score vectors and tweet content to detect logical contradictions.



Model Optimization: Surpass current metrics by training a BERT model.



Data Augmentation: Expand the dataset through web scraping of tweets.